

Tracking the Acceleration: An Empirical Audit of Dario Amodei's AI Timeline Predictions

Report generated 2026-06-16. All claims are linked to explicit primary sources. This is an independent analytical project, not affiliated with Anthropic.

Abstract

We maintain a continuously-updated, primary-source-anchored database of public predictions made by Dario Amodei, CEO of Anthropic, concerning the arrival and impact of "powerful AI." As of this revision the corpus contains 6 tracked predictions (Pending: 4, Failed: 1, Walked Back: 1). We test one principal hypothesis (H1): that Amodei's successive "powerful AI" timelines exhibit a *contracting prediction horizon* ? the interval between when a forecast is made and its target date shrinks over calendar time ? which is the observable signature of an accelerating (super-exponential) capability expectation. A linear fit yields a horizon slope of -0.1974 years per calendar year ($R^2 = 0.9659$), supported of H1. We contextualise this with a first-order effective-compute growth model and present the full prediction timeline as the report's central artifact.

Introduction

Anthropic's public posture rests on a *scaling paradigm*: the thesis that continued growth in training compute, data, and algorithmic efficiency yields broadly predictable gains in model capability, extrapolating toward systems Amodei describes as a "country of geniuses in a datacenter." Beginning with the October 2024 essay *Machines of Loving Grace* and continuing through the January 2026 essay *The Adolescence of Technology*, Amodei has issued a series of dated, falsifiable forecasts spanning AGI timelines, software automation, biomedical acceleration, and labour-market disruption.

These forecasts are unusually concrete for a frontier-lab CEO, which makes them auditable. This tracker exists to hold each claim against its primary source and against unfolding events, and ? critically ? to quantify whether the *cadence* of the forecasts is itself accelerating.

Methodology (How the Science Was Done)

Data gathering. Each prediction is logged in 'Mem/predictions_db.json' only when an explicit primary source (essay, policy submission, recorded interview, or directly-quoted public remarks) can be hyperlinked. Statements were located via targeted web search and cross-checked against the canonical publication (e.g. 'darioamodei.com' for essays). Entries carry the verbatim quote where available, the date made, an inferred target date, a status, and an assessment.

Status verification. Each prediction is assigned one of: Pending, Achieved, Partial, Failed, or Walked Back. A claim moves to Failed when its explicit target window elapses without realization (e.g. the "90% of code in 3?6 months" claim, whose September 2025 window passed), and to Walked Back when the predictor publicly softens or retracts it (corroborated by reporting).

Thesis testing. 'Scripts/thesis_acceleration.py' computes, for every AGI-timeline prediction, the horizon = target ? made, then regresses horizon on date-made. A significantly negative slope supports H1. A supporting effective-compute model projects relative compute growth at an assumed 2.5 doublings/year (? 5.66×annually) to contextualise the ~2027 window. All code is pure-standard-library and reproducible via 'Scripts/run_pipeline.sh'.

Timeline Analysis (Primary Function)

The table below maps every tracked prediction from issuance to target, with its realized horizon and current status. This timeline is the core artifact of the project; the interactive version lives on the GitHub Pages site ('docs/').

ID	Category	Made	Target	Horizon (yr)	Status	Source
P001	Powerful AI	2024-10-24	2026-12-31	2.18	Pending	Machines of Loving Grace (https://darioamodei.com/essay/machines-of-loving-grace)
P005	Biology & Health	2025-01-21	2035-12-31	10.94	Pending	World Economic Forum (Davos) 2025 remarks (reported) (https://www.pymnts.com/artificial-intelligence-2/2025/anthropic-ceo-sees-ai-powered-advances-doubling-human-lifespans/)
P002	Powerful AI	2025-03-06	2027-03-31	2.07	Pending	Anthropic Recommendations to OSTP for the U.S. AI Action Plan (https://www.anthropic.com/news/anthropic-s-recommendations-ostp-u-s-ai-action-plan)

| P004 | Software & Code Automation | 2025-03-10 | 2025-09-10 | 0.50 | Failed | Council on Foreign Relations event (reported)
(<https://finance.yahoo.com/news/anthropic-ceo-says-ai-could-193020957.html>) |

| P006 | Labor & Economy | 2025-05-28 | 2030-05-28 | 5.00 | Walked Back | Axios interview: 'Behind the Curtain: A white-collar bloodbath'
(<https://www.axios.com/2025/05/28/ai-jobs-white-collar-unemployment-anthropic>) |

| P003 | Powerful AI | 2026-01-27 | 2027-12-31 | 1.92 | Pending | The Adolescence of Technology (<https://darioamodei.com/essay/the-adolescence-of-technology>) |

H1 ? contracting horizon. Regressing horizon on date-made across the
AGI-timeline predictions gives slope = -0.1974 yr/yr

($R^2 = 0.9659$). SUPPORTED: prediction horizon is contracting over time (-0.20 yr of horizon per calendar yr), the signature of an accelerating (super-exponential) capability expectation.

Supporting context ? effective compute. Under the stated assumption, relative effective compute grows as:

| Year | Relative effective compute (x 2024) |

| 2024 | 1.0x |

| 2025 | 5.7x |

| 2026 | 32.0x |

| 2027 | 181.0x |

| 2028 | 1,024.0x |

| 2029 | 5,792.6x |

| 2030 | 32,768.0x |

This projection is illustrative, not a claim of Amodei's; it situates why a 2026?2027 "powerful AI" window is internally consistent with the scaling paradigm even as specific sub-claims (e.g. code automation) have slipped.

Per-prediction detail

P001 ? Powerful AI / AGI Timeline

Claim. "Powerful AI" (a 'country of geniuses in a datacenter') could arrive as early as 2026 in the optimistic scenario.

"I think it could come as early as 2026... a country of geniuses in a datacenter."

Made: 2024-10-24 • Target: 2026-12-31 • Status: Pending

Primary source: Machines of Loving Grace

(<https://darioamodei.com/essay/machines-of-loving-grace>) (essay, darioamodei.com)

Assessment. Defines 'powerful AI' as smarter than a Nobel laureate across most fields, with the agency, tools, and speed of millions of copies. Seed entry.

P005 ? Biology & Health

Claim. AI could compress 50-100 years of biological progress into 5-10 years, potentially doubling the human lifespan.

"I think AI could help us compress the next 50 to 100 years of biological progress into 5 to 10 years."

Made: 2025-01-21 • Target: 2035-12-31 • Status: Pending

Primary source: World Economic Forum (Davos) 2025 remarks (reported) (<https://www.pymnts.com/artificial-intelligence-2/2025/anthropic-ceo-sees-ai-powered-advances-doubling-human-lifespans/>) (interview, World Economic Forum / PYMNTS)

Assessment. The 'compressed 21st century' thesis. Reasoning: 20th-century life expectancy roughly doubled (~40 to ~75); a compressed 21st could double it again toward ~150. Long-horizon; remains Pending.

P002 ? Powerful AI / AGI Timeline

Claim. Powerful AI systems will emerge in late 2026 or early 2027, with intellectual capabilities matching or exceeding Nobel Prize winners across most disciplines.

"we expect powerful AI systems will emerge in late 2026 or early 2027... matching or exceeding that of Nobel Prize winners across most disciplines."

Made: 2025-03-06 • Target: 2027-03-31 • Status: Pending

Primary source: Anthropic Recommendations to OSTP for the U.S. AI Action Plan (<https://www.anthropic.com/news/anthropic-s-recommendations-ostp-u-s-ai-action-plan>) (policy submission, Anthropic)

Assessment. Formal written policy submission; the most concrete dated AGI-timeline commitment on record.

P004 ? Software & Code Automation

Claim. AI will be writing 90% of all code within 3-6 months, and essentially all code within 12 months.

"I think we'll be there in three to six months, where AI is writing 90 percent of the code. And then, in 12 months, we may be in a world where AI is writing essentially all of the code."

Made: 2025-03-10 • Target: 2025-09-10 • Status: Failed

Primary source: Council on Foreign Relations event (reported) (<https://finance.yahoo.com/news/anthropic-ceo-says-ai-could-193020957.html>) (public remarks, Council on Foreign Relations / Yahoo Finance)

Assessment. The 3-6 month window (Sep 2025) passed without the 90%-of-all-code threshold being met industry-wide. Amodei later softened to 'on many teams.' Independent analyses (Redwood Research) found the claim unverified even within Anthropic.

P006 ? Labor & Economy

Claim. AI could eliminate up to 50% of entry-level white-collar jobs within five years and

push unemployment to 10-20%.

"AI could wipe out half of all entry-level white-collar jobs and spike unemployment to 10-20% in the next one to five years."

Made: 2025-05-28 • Target: 2030-05-28 • Status: Walked Back

Primary source: Axios interview: 'Behind the Curtain: A white-collar bloodbath' (<https://www.axios.com/2025/05/28/ai-jobs-white-collar-unemployment-anthropic>) (interview, Axios)

Assessment. By May 2026, Amodei publicly tempered this forecast (Fortune), suggesting automation may expand work rather than purely destroy it. Yale Budget Lab found no meaningful unemployment shift for AI-exposed workers through early 2026, though entry-level tech hiring did fall sharply.

P003 ? Powerful AI / AGI Timeline

Claim. A 'country of geniuses in a datacenter' (~50 million AI agents, each smarter than a Nobel laureate, running at 10-100x human speed) could materialize by 2027.

"a country of geniuses in a datacenter"

Made: 2026-01-27 • Target: 2027-12-31 • Status: Pending

Primary source: The Adolescence of Technology (<https://darioamodei.com/essay/the-adolescence-of-technology>) (essay, darioamodei.com)

Assessment. Companion essay to Machines of Loving Grace, published Jan 2026; reaffirms the ~2027 powerful-AI window while emphasizing five risk categories.

Conclusions

The evidence base is mixed but structured. On AGI timing, the load-bearing forecasts (P002, P003) remain **Pending** and squarely live: the late-2026/early-2027 window has not yet elapsed, and the cadence of Amodei's forecasts is measurably **tightening** (H1 supported), consistent with an accelerationist read of the scaling curve. On near-term operational claims, the record is weaker: the "90% of code" forecast (P004) failed its explicit window, and the "white-collar bloodbath" labour forecast (P006) has been walked back amid employment data that has so far not confirmed it.

The synthesis: Amodei is more calibrated on **capability emergence at the research frontier** than on **the speed of real-world diffusion** (code, jobs), where institutional and adoption frictions dominate. The next revisions of this report will watch the 2026?2027 window closely, as it is decisive for P001?P003.

**Reproduce: 'bash Scripts/run_pipeline.sh'. Data and sources: 'Mem/'. Generated by an*

autonomous R&D agent.*